

Algorithms for Data Science

CSOR W4246

Eleni Drinea
Computer Science Department

Columbia University

Tuesday, September 8, 2015

Outline

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms
- 4 Efficient algorithms
- 5 Asymptotic notation

Today

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms
- 4 Efficient algorithms
- 5 Asymptotic notation

Algorithms

- ▶ An **algorithm** is a well-defined computational procedure that transforms the **input** (a set of values) into the **output** (a new set of values).
- ▶ The desired input/output relationship is specified by the statement of the **computational problem** for which the algorithm is designed.
- ▶ An algorithm is **correct** if, for every input, it **halts** with the correct output.

Efficient Algorithms

- ▶ In this course we are interested in algorithms that are **correct** and **efficient**.
- ▶ Efficiency is related to the **resources** an algorithm uses:
time, space
 - ▶ *How much time/space are used?*
 - ▶ *How do they **scale** as the input size grows?*

We will primarily focus on efficiency in **running time**.

Running time

Running time = number of **primitive computational steps** performed; typically these are

1. arithmetic operations: add, subtract, multiply, divide **fixed-size** integers
2. data movement operations: load, store, copy
3. control operations: branching, subroutine call and return

We will use **pseudocode** for our algorithm descriptions.

Today

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms
- 4 Efficient algorithms
- 5 Asymptotic notation

Sorting

- ▶ **Input:** A list A of n integers x_1, \dots, x_n .
- ▶ **Output:** A permutation x'_1, x'_2, \dots, x'_n of the n integers where they are sorted in non-decreasing order, i.e.,
$$x'_1 \leq x'_2 \leq \dots \leq x'_n$$

Sorting

- ▶ **Input:** A list A of n integers x_1, \dots, x_n .
- ▶ **Output:** A permutation x'_1, x'_2, \dots, x'_n of the n integers where they are sorted in non-decreasing order, i.e.,
$$x'_1 \leq x'_2 \leq \dots \leq x'_n$$

Example

- ▶ Input: $n = 6$, $A = \{9, 3, 2, 6, 8, 5\}$
- ▶ Output: $A = \{2, 3, 5, 6, 8, 9\}$

What *data structure* should we use to represent the list?

Sorting

- ▶ **Input:** A list A of n integers x_1, \dots, x_n .
- ▶ **Output:** A permutation x'_1, x'_2, \dots, x'_n of the n integers where they are sorted in non-decreasing order, i.e.,
$$x'_1 \leq x'_2 \leq \dots \leq x'_n$$

Example


- ▶ Input: $n = 6$, $A = \{9, 3, 2, 6, 8, 5\}$
- ▶ Output: $A = \{2, 3, 5, 6, 8, 9\}$

What *data structure* should we use to represent the list?

Array: collection of items of the same data type


- ▶ allows for *random access*
- ▶ “zero” indexed in C++ and Java

Insertion sort in English




1. Start with a (trivially) sorted subarray of size 1 consisting of the first element of A .

Insertion sort in English



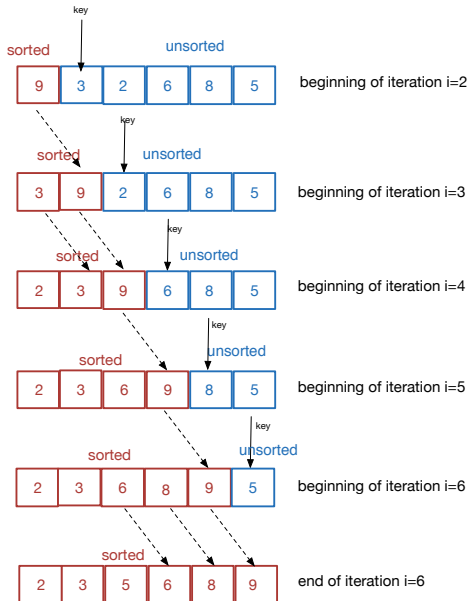
1. Start with a (trivially) sorted subarray of size 1 consisting of the first element of A .
2. **Increase** the size of the sorted subarray by 1 by **inserting** the next element of A into its **correct** location.
 - ▶ Compare that next element, call it **key**, with every element x of the sorted subarray starting from the **right**.
 - ▶ If $x > \text{key}$, move x one position to the right.
 - ▶ Else ($x \leq \text{key}$), **insert** **key** after x .

Insertion sort in English



1. Start with a (trivially) sorted subarray of size 1 consisting of the first element of A .
2. **Increase** the size of the sorted subarray by 1 by **inserting** the next element of A into its **correct** location.
 - ▶ Compare that next element, call it **key**, with every element x of the sorted subarray starting from the **right**.
 - ▶ If $x > \text{key}$, move x one position to the right.
 - ▶ Else ($x \leq \text{key}$), **insert** **key** after x .
3. Repeat Step 2. until the sorted subarray has size n .

Example of insertion sort: $n = 6, A = \{9, 3, 2, 6, 8, 5\}$



Pseudocode

Let A be an array of n integers.

insertion-sort(A)

for $i = 2$ to n **do**

$\text{key} = A[i]$

 //Insert $A[i]$ into the sorted subarray $A[1, i - 1]$

$j = i - 1$

while $j > 0$ and $A[j] > \text{key}$ **do**

$A[j + 1] = A[j]$

$j = j - 1$

end while

$A[j + 1] = \text{key}$

end for

Today

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms**
- 4 Efficient algorithms
- 5 Asymptotic notation

Analysis of algorithms

- ▶ **Correctness**
- ▶ **Running time**
- ▶ **(Space)**

Analysis of algorithms

- ▶ **Correctness:** formal proof often by **induction**
- ▶ **Running time:** number of **primitive computational steps**
 - ▶ Not the same as **time** it takes to execute the algorithm.
 - ▶ We want a measure that is independent of hardware.
 - ▶ We want to know how running time **scales** with the size of the input.
- ▶ **Space:** how much space is required by the algorithm

Analysis of insertion-sort

Notation: $A[i, j]$ is the subarray of A that starts at position i and ends at position j .

- ▶ **Correctness:** follows from the key observation that after loop i , the subarray $A[1, i]$ is sorted
- ▶ **Running time:** number of primitive computational steps
- ▶ **Space:** **in place algorithm** (at most a constant number of elements of A are stored outside A at any time)

Example of induction

Fact 1.

For all $n \geq 1$, $\sum_{i=1}^n i = \frac{n(n+1)}{2}$.

Example of induction

Fact 1.

For all $n \geq 1$, $\sum_{i=1}^n i = \frac{n(n+1)}{2}$.

Proof.

- ▶ **Base case:** $n = 1$
- ▶ **Inductive hypothesis:** Assume that the statement is true for $n \geq 1$, that is, $\sum_{i=1}^n i = \frac{n(n+1)}{2}$.
- ▶ **Inductive step:** We show that the statement is true for $n + 1$. That is, $\sum_{i=1}^{n+1} i = \frac{(n+1)(n+2)}{2}$. (Show this!)
- ▶ **Conclusion:** It follows that the statement is true for all n since we can apply the inductive step for $n = 2, 3, \dots$



Correctness of insertion-sort

Notation: $A[i, j]$ is the subarray of A that starts at position i and ends at position j .

Minor change in the pseudocode: in line 1, start from $i = 1$ rather than $i = 2$. *How does this change affect the algorithm?*

Claim 1.

Let $n \geq 1$ be a positive integer. For all $1 \leq i \leq n$, after the i -th loop, the subarray $A[1, i]$ is sorted.

Correctness of `insertion-sort` follows if we show Claim 1 (*why?*).

Proof of Claim 1

By induction on i .

- ▶ **Base case:** $i = 1$, trivial.
- ▶ **Induction hypothesis:** assume that the statement is true for some $1 \leq i < n$.
- ▶ **Inductive step:** Show it true for $i + 1$.

In loop $i + 1$, element $\text{key} = A[i + 1]$ is inserted into $A[1, i]$, which is sorted (by the induction hypothesis). Since key is inserted after the first element $A[\ell]$ for $1 \leq \ell \leq i$ such that $\text{key} \geq A[\ell]$, and all elements in $A[\ell + 1, j]$ are pushed one position to the right with their order preserved, the statement is true for $i + 1$.

Running time $T(n)$ of insertion-sort

```
for  $i = 2$  to  $n$  do  
     $\text{key} = A[i]$   
    //Insert  $A[i]$  into the sorted subarray  $A[1, i - 1]$   
     $j = i - 1$   
    while  $j > 0$  and  $A[j] > \text{key}$  do  
         $A[j + 1] = A[j]$   
         $j = j - 1$   
    end while  
     $A[j + 1] = \text{key}$   
end for
```

- ▶ How many *primitive computational steps* are executed by the algorithm?
- ▶ Equivalently, what is the running time $T(n)$? Bounds on $T(n)$?

Running time $T(n)$ of insertion-sort

```
for  $i = 2$  to  $n$  do                                line 1
    key =  $A[i]$                                        line 2
    //Insert  $A[i]$  into the sorted subarray  $A[1, i - 1]$ 
     $j = i - 1$                                        line 3
    while  $j > 0$  and  $A[j] > \mathbf{key}$  do             line 4
         $A[j + 1] = A[j]$                              line 5
         $j = j - 1$                                    line 6
    end while
     $A[j + 1] = \mathbf{key}$                              line 7
end for
```

- For $2 \leq i \leq n$, let $t_i = \#$ times line 4 is executed.

Running time $T(n)$ of insertion-sort

```
for  $i = 2$  to  $n$  do           line 1
    key =  $A[i]$                 line 2
    //Insert  $A[i]$  into the sorted subarray  $A[1, i - 1]$ 
     $j = i - 1$                 line 3
    while  $j > 0$  and  $A[j] > \text{key}$  do line 4
         $A[j + 1] = A[j]$       line 5
         $j = j - 1$            line 6
    end while
     $A[j + 1] = \text{key}$         line 7
end for
```

- For $2 \leq i \leq n$, let $t_i = \#$ times line 4 is executed. Then

$$T(n) = n + 3(n - 1) + \sum_{i=2}^n t_i + 2 \sum_{i=2}^n (t_i - 1) = 3 \sum_{i=2}^n t_i + 2n - 1$$

- Which input yields the smallest (best-case) running time?
- Which input yields the largest (worst-case) running time?

Running time $T(n)$ of insertion-sort

```
for  $i = 2$  to  $n$  do                                line 1
     $\text{key} = A[i]$                                     line 2
    //Insert  $A[i]$  into the sorted subarray  $A[1, i - 1]$ 
     $j = i - 1$                                         line 3
    while  $j > 0$  and  $A[j] > \text{key}$  do              line 4
         $A[j + 1] = A[j]$                             line 5
         $j = j - 1$                                     line 6
    end while
     $A[j + 1] = \text{key}$                                 line 7
end for
```

- For $2 \leq i \leq n$, let $t_i = \#$ times line 4 is executed. Then

$$T(n) = 3 \sum_{i=2}^n t_i + 2n - 1$$

- **Best-case** running time: $5n - 4$
- **Worst-case** running time: $\frac{3n^2}{2} + \frac{7n}{2} - 4$

Definition 2.

Worst-case running time: largest possible running time of the algorithm over all inputs of a given size n .

Why *worst-case* analysis?

- ▶ It gives well-defined computable bounds.
- ▶ Average-case analysis can be tricky: how do we generate a “random” instance?

The worst-case running time of `insertion-sort` is quadratic.
Is `insertion-sort` *efficient*?

Today

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms
- 4 Efficient algorithms**
- 5 Asymptotic notation

Efficiency of insertion-sort and the brute force solution

Compare to **brute force** solution:

- ▶ At each step, generate a new permutation of the n integers.
- ▶ If sorted, stop and output the permutation.

Efficiency of insertion-sort and the brute force solution

Compare to **brute force** solution:

- ▶ At each step, generate a new permutation of the n integers.
- ▶ If sorted, stop and output the permutation.

Worst-case analysis: generate $n!$ permutations. *Is brute force solution efficient?*

Efficiency of insertion-sort and the brute force solution

Compare to **brute force** solution:

- ▶ At each step, generate a new permutation of the n integers.
- ▶ If sorted, stop and output the permutation.

Worst-case analysis: generate $n!$ permutations. *Is brute force solution efficient?*

- ▶ Efficiency relates to the performance of the algorithm as n grows.
- ▶ Stirling's approximation formula: $n! \approx \left(\frac{n}{e}\right)^n$.
 - ▶ For $n = 10$, generate $3.67^{10} \geq 2^{10}$ permutations.
 - ▶ For $n = 50$, generate $18.3^{50} \geq 2^{200}$ permutations.
 - ▶ For $n = 100$, generate $36.7^{100} \geq 2^{700}$ permutations!

⇒ Brute force solution is **not** efficient.

Definition 3 (Attempt 1).

An algorithm is efficient if it achieves better worst-case performance than brute-force search.

Definition 3 (Attempt 1).

An algorithm is efficient if it achieves better worst-case performance than brute-force search.

Caveat: fails to discuss the **scaling properties** of the algorithm.

- ▶ If the input size grows by a constant factor, we would like the running time $T(n)$ of the algorithm to increase by a constant factor as well.

Definition 3 (Attempt 1).

An algorithm is efficient if it achieves better worst-case performance than brute-force search.

Caveat: fails to discuss the **scaling properties** of the algorithm.

- ▶ If the input size grows by a constant factor, we would like the running time $T(n)$ of the algorithm to increase by a constant factor as well.
- ▶ Note that **polynomial running times scale well**: on input of size n , $T(n)$ is at most $c \cdot n^d$ for $c, d > 0$ constants.
 - ▶ the **smaller** the exponent of the polynomial the better

Definition 4.

An algorithm is efficient if it has a polynomial running time.

Caveat

- ▶ What about huge constants in front of the leading term or large exponents?

However

- ▶ **Small degree polynomial** running times exist for most problems that can be solved in polynomial time.
- ▶ Conversely, problems for which no polynomial-time algorithm is known tend to be very hard in practice.
- ▶ So we can distinguish between **easy** and **hard** problems.

Remark 1.

Today's big data: even low degree polynomials might be too slow!

Are we done with sorting?

Insertion sort is efficient. *Are we done with sorting?*

Are we done with sorting?

Insertion sort is efficient. *Are we done with sorting?*

1. *Can we do better?*
2. *And what is better?*
 - ▶ *E.g., is $T(n) = n^2 + n - 4$ worth aiming for?*

Running time in terms of # primitive steps

To discuss this, we need a coarser classification of running times of algorithms; exact characterizations

- ▶ are **too detailed**;
- ▶ do not reveal similarities between running times in an immediate way as n grows large;
- ▶ are often **meaningless**: pseudocode steps will **expand** by a constant factor that depends on the hardware.

Asymptotic analysis

A framework that will allow us to compare the **rate of growth** of different running times as the input size n grows.

- ▶ We will express the running time as a function of the number of primitive steps.
 - ▶ The number of primitive steps is itself a function of the size of the input n .
- ⇒ The running time is a function of the size of the input n .
- ▶ To compare functions expressing running times, **we will ignore their low-order terms and focus solely on the highest-order term.**


Today

- 1 Overview
- 2 A first algorithm: insertion sort
- 3 Analysis of algorithms
- 4 Efficient algorithms
- 5 Asymptotic notation**

Asymptotic upper bounds: Big- O notation

Definition 5 (O).

We say that $T(n) = O(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) \leq c \cdot f(n)$.



Asymptotic upper bounds: Big- O notation

Definition 6 (O).

We say that $T(n) = O(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) \leq c \cdot f(n)$.


Examples:

- ▶ $T(n) = an^2 + b$, $a, b > 0$ constants and $f(n) = n^2$.
- ▶ $T(n) = an^2 + b$, $f(n) = n^3$.

Asymptotic lower bounds: Big- Ω notation

Definition 7 (Ω).

We say that $T(n) = \Omega(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) \geq c \cdot f(n)$.



Asymptotic lower bounds: Big- Ω notation

Definition 8 (Ω).

We say that $T(n) = \Omega(f(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) \geq c \cdot f(n)$.

Examples:


- ▶ $T(n) = an^2 + b$, $a, b > 0$ constants and $f(n) = n^2$.
- ▶ $T(n) = an^2 + b$, $a, b > 0$ constants and $f(n) = n$.

Asymptotic tight bounds: Θ notation

Definition 9 (Θ).

We say that $T(n) = \Theta(f(n))$ if there exist constants $c_1, c_2 > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have

$$c_1 \cdot f(n) \leq T(n) \leq c_2 \cdot f(n).$$



Asymptotic tight bounds: Θ notation

Definition 10 (Θ).

We say that $T(n) = \Theta(f(n))$ if there exist constants $c_1, c_2 > 0$ and $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have

$$c_1 \cdot f(n) \leq T(n) \leq c_2 \cdot f(n).$$

Equivalent definition

$T(n) = \Theta(f(n))$ if $T(n) = O(f(n))$ and $T(n) = \Omega(f(n))$

Examples:

- ▶ $T(n) = an^2 + b$, $a, b > 0$ constants and $f(n) = n^2$.
- ▶ $T(n) = n \log n + n$, and $f(n) = n \log n$.

Asymptotic upper bounds that are **not** tight: little- o

Definition 11 (o).

We say that $T(n) = o(f(n))$ if **for any** constant $c > 0$ there exists a constant $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) < c \cdot f(n)$.

Asymptotic upper bounds that are **not** tight: little- o

Definition 11 (o).

We say that $T(n) = o(f(n))$ if **for any** constant $c > 0$ there exists a constant $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) < c \cdot f(n)$.

- ▶ Intuitively, $T(n)$ becomes **insignificant** relative to $f(n)$ as $n \rightarrow \infty$.
- ▶ Proof by showing that $\lim_{n \rightarrow \infty} \frac{T(n)}{f(n)} = 0$ (if the limit exists).

Asymptotic upper bounds that are **not** tight: little- o

Definition 11 (o).

We say that $T(n) = o(f(n))$ if **for any** constant $c > 0$ there exists a constant $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) < c \cdot f(n)$.

- ▶ Intuitively, $T(n)$ becomes **insignificant** relative to $f(n)$ as $n \rightarrow \infty$.
- ▶ Proof by showing that $\lim_{n \rightarrow \infty} \frac{T(n)}{f(n)} = 0$ (if the limit exists).

Examples:

- ▶ $T(n) = an^2 + b$, $a, b > 0$ constants and $f(n) = n^3$.
- ▶ $T(n) = n \log n$, $a, b, d > 0$ constants and $f(n) = n^2$.

Asymptotic lower bounds that are **not** tight: little- ω

Definition 12 (ω).

We say that $T(n) = \omega(f(n))$ if **for any** constant $c > 0$ there exists $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) > c \cdot f(n)$.

Asymptotic lower bounds that are **not** tight: little- ω

Definition 12 (ω).

We say that $T(n) = \omega(f(n))$ if **for any** constant $c > 0$ there exists $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) > c \cdot f(n)$.

- ▶ Intuitively $T(n)$ becomes **arbitrarily large** relative to $f(n)$ as $n \rightarrow \infty$.
- ▶ $T(n) = \omega(f(n))$ implies that $\lim_{n \rightarrow \infty} \frac{T(n)}{f(n)} = \infty$ if the limit exists. Then $f(n) = o(T(n))$.

Asymptotic lower bounds that are **not** tight: little- ω

Definition 12 (ω).

We say that $T(n) = \omega(f(n))$ if **for any** constant $c > 0$ there exists $n_0 \geq 0$ s.t. for all $n \geq n_0$, we have $T(n) > c \cdot f(n)$.

- ▶ Intuitively $T(n)$ becomes **arbitrarily large** relative to $f(n)$ as $n \rightarrow \infty$.
- ▶ $T(n) = \omega(f(n))$ implies that $\lim_{n \rightarrow \infty} \frac{T(n)}{f(n)} = \infty$ if the limit exists. Then $f(n) = o(T(n))$.

Examples:

- ▶ $T(n) = n^2$ and $f(n) = n \log n$.
- ▶ $T(n) = 2^n$ and $f(n) = n^5$.

Basic rules for omitting low order terms from functions

1. ignore **multiplicative** factors: e.g., $10n^3$ becomes n^3
 2. n^a dominates n^b if $a > b$: e.g., n^2 dominates n
 3. exponentials dominate polynomials: e.g., 2^n dominates n^4
 4. polynomials dominate logarithms: e.g., n dominates $\log^3 n$
- \Rightarrow for large enough n ,

$$\log n < n < n \log n < n^2 < 2^n < n^n$$

Properties of asymptotic growth rates

Transitivity

1. If $f = O(g)$ and $g = O(h)$ then $f = O(h)$.
2. If $f = \Omega(g)$ and $g = \Omega(h)$ then $f = \Omega(h)$.
3. If $f = \Theta(g)$ and $g = \Theta(h)$ then $f = \Theta(h)$.

Sums of (up to a constant number of) functions

1. If $f = O(h)$ and $g = O(h)$ then $f + g = O(h)$.
2. Let k be a fixed constant, and let f_1, f_2, \dots, f_k, h be functions s.t. for all i , $f_i = O(h)$. Then $f_1 + f_2 + \dots + f_k = O(h)$.

Transpose symmetry

- ▶ $f = O(g)$ if and only if $g = \Omega(f)$.
- ▶ $f = o(g)$ if and only if $g = \omega(f)$.